


FitPredict

By: Havyn Nguyen, Arpita Pandey, Prakhar Shah





01

Predictive
Task:

INTRODUCTION

Our goal is to predict the star rating (1–5) that a user will give to a clothing item based on past reviews and item attributes. This is a personalized recommendation problem formulated as regression. A strong model should account for differences in both individual users and specific clothing items, as our EDA showed large variation in preferences and item popularity.



Evaluation Strategy

We evaluate performance using RMSE and MAE, which measure how far predictions deviate from true ratings. These metrics are appropriate because our data is heavily skewed toward high ratings, making accuracy-based metrics misleading. We also report performance on a held-out test split to measure generalization.



Baselines for Comparison

Global Mean Baseline

always predict the average rating




Bias Model

global mean + user and item biases
(accounts for harsher reviewers and more-liked items)



These baselines establish whether more complex models are actually providing value.

Validity?



We assess whether predictions make sense both quantitatively and qualitatively:

- Improved RMSE/MAE vs. baselines → better performance
- Per-group error (e.g., fit labels) → fairness and robustness checks
 - Rating distributions → ensure predicted values match real behavior

This combination ensures our model not only minimizes error, but also reflects real consumer experiences.





02 Exploratory Data Analysis

Dataset Context: Clothing Fit & Review Data (Rent the Runway)

- Collected from a real clothing rental platform
- Contains user reviews, ratings, fit labels, and body measurements
- Goal: Understand how fit + body type influence satisfaction
- Supports personalized size recommendations

Includes ~192K reviews from ~100K users and ~6K clothing items




Data Preprocessing

Data Cleaning & Feature Processing

- Converted star ratings from 2–10 → 1–5 scale
- Parsed height and weight into numeric values
- Filtered dataset to valid numeric user + rating data
- Kept 3 main fit labels: fit, small, large

	fit	user_id	bust size	item_id	weight	rating	rented for	review_text	body type	review_summary	category	height	size	age	review_date	stars
0	fit	420272	34d	2280466	137lbs	10	vacation	An adorable romper! Belt and zipper were a lit...	hourglass	So many compliments!	romper	5' 8"	14	28	April 20, 2016	5.0
1	fit	273551	34b	153475	132lbs	10	other	I rented this dress for a photo shoot. The the...	straight & narrow	I felt so glamorous!!!	gown	5' 8"	12	36	June 18, 2013	5.0
2	fit	360448	NaN	1063761	NaN	10	party	This hugged in all the right places! It was a ...	NaN	It was a great time to celebrate the (almost) ...	sheath	5' 4"	4	116	December 14, 2015	5.0
3	fit	909926	34c	126335	135lbs	8	formal affair	I rented this for my company's black tie award...	pear	Dress arrived on time and in perfect condition.	dress	5' 5"	8	34	February 12, 2014	4.0
4	fit	151944	34b	616682	145lbs	10	wedding	I have always been petite in my upper body and...	athletic	Was in love with this dress !!!	gown	5' 9"	12	27	September 26, 2016	5.0

Strong Positivity Bias in our Data



This plot shows a heavily skewed rating distribution, with 5-star reviews dominating (~65%) and <3-star ratings extremely rare.


This indicates a strong positivity bias typical in consumer review datasets.

Modeling implications:

- Low variance in labels reduces the discriminative power of simple models
- A trivial predictor (always predict 5★) would score high accuracy
- Therefore, RMSE/MAE are more appropriate evaluation metrics than accuracy
- Predictive models must capture subtle preference differences to be meaningful

This motivates the need for personalized recommendation approaches rather than naïve baselines.

Fit Label Distribution




Most reviews indicate the item fit true-to-size (~74%), while only about 13% report “small” and 13% report “large”. This confirms that users mainly review when satisfied with fit, but there remains a significant portion of feedback describing sizing issues. From a modeling perspective:

- The fit label is not uniformly distributed, so models must learn from an imbalanced categorical feature
- Fit labels provide high signal for rating prediction, as seen in later analysis
- Including this feature enables better personalization around size suitability

The distribution reinforces that fit information is a valuable predictive input that can explain variation in user satisfaction beyond star ratings alone.

Fit Feedback Correlates Strongly with Satisfaction

Fit Feedback Strongly Predicts Rating



The boxplot shows a clear rating gradient based on fit labels:


- True-to-size reviews average ~4.65★
- “Small” fit reviews drop to ~4.21★
- “Large” fit reviews drop to ~4.28★

Even though most reviews are positive overall, size dissatisfaction consistently lowers ratings. This highlights two important modeling implications:

- Fit feedback is a high-value predictive feature for rating prediction
- Users penalize sizing inaccuracies more than other product qualities

This insight motivates including fit label (and potentially body measurements) in our feature set to better capture user-item preference interactions.

Body Measurements



Height and weight follow realistic, approximately normal distributions:


- Mean height ~65 in (5'5")
- Mean weight ~140 lb
- Most users fall within typical adult female ranges

This confirms:

- The dataset reflects a real consumer population, useful for modeling personal fit
- These features provide continuous variability that can improve rating predictions beyond categorical fit labels
- Outliers are visible but not extreme → we can retain most values without heavy filtering

Since body measurements directly influence how clothing fits, incorporating height and weight helps the model personalize predictions to each reviewer's body type.

Item Popularity — Extreme Long-Tail Distribution



The log-scaled histogram shows a heavy long-tail:


- Majority of items receive very few reviews (sparsity)
- A small number of popular items dominate the data
- ~80% of items have <10 reviews, while a few exceed 1000+

This aligns with typical recommender datasets where exposure is uneven. For modeling, this implies:

- We must handle cold-start items with limited observations
- Matrix Factorization should still work well for popular items, while content features (e.g., category, measurements) may help generalize to less-reviewed items

Understanding item sparsity directly informs how we design and evaluate our predictive models.

Most Users Contribute Very Few Reviews



This bar chart highlights extreme user-side sparsity:

- Over 68% of users write only 1 review
- The number of multi-review users drops exponentially
- Only a tiny fraction produce enough data to learn reliable user embeddings

This reinforces a key modeling challenge:

Collaborative filtering alone will underperform.

We must leverage item metadata and review text to make robust predictions, especially for new or low-activity users.

Users with exactly 1 review: 68.02517344656329 %

Users with ≥ 5 reviews: 5.428972210638055 %

Max reviews by a single user: 436

Textual Data and Preprocessing

Our dataset provided two textual fields for each review: a short “**review_summary**” and a longer “**review_text**”. We chose to combine these into a single unified column called “**review**” because the summary often repeats key descriptive phrases and merging both fields provides the model with more complete linguistic context. While the summary provides a brief, high-level description and the full text offers detailed context, treating them separately can dilute the signal available to the model. Merging them ensures that all textual information associated with a review is captured in one place, simplifies preprocessing, and allows downstream feature extraction methods to operate on a richer, more coherent input. This consolidated text field ultimately strengthens the quality of the learned representations and improves the expressiveness of any text-based modeling.

The resulting review column contains no missing entries in review field, confirming that every observation includes a complete written review.

Before feature extraction, we applied a full preprocessing pipeline that included lowercasing, punctuation removal, tokenization, and stemming using the Porter stemmer. These steps reduced noise, standardized vocabulary, and helped eliminate redundant word variations.



Textual Feature Engineering

We evaluated three main text-representation strategies: **Bag-of-Words with multi-length n-grams, TF-IDF, and Word2Vec embeddings**. The Bag-of-Words model counted occurrences of unigrams through five-grams, producing a very high-dimensional sparse matrix. TF-IDF used the same tokenization structure but reweighted terms based on their distinctiveness across the corpus, highlighting rare but informative descriptors while down-weighting generic words. Word2Vec created 100-dimensional dense embeddings by training a shallow neural model on the collection of reviews, after which each review was represented by the average of its word vectors. These three approaches provided complementary perspectives: raw frequency, importance-weighted frequency, and semantic similarity.

Textual Feature Engineering

Text Feature	Computational Time	Why Used	Challenges
Bag-of-Words	~2 hours	Simple, interpretable	Huge sparse matrix → RAM issues
TF-IDF	~15 minutes	Captures important keywords	High dimensionality
Word2Vec embeddings	~15 minutes	Semantic similarity, smaller vectors	Requires careful preprocessing



Textual Feature Engineering

To evaluate the effectiveness of our text-feature representations, we trained separate predictive models for each method—Bag-of-Words, TF-IDF, and Word2Vec—and computed three key metrics: **MSE, RMSE, and AUC**.

For regression performance, we used Ridge Regression, which is well-suited for high-dimensional sparse text data, including tens of thousands of n-grams and TF-IDF features. Ridge Regression minimizes squared error while penalizing large weights, making it an appropriate choice for text models with many correlated features. For AUC, we reframed the task temporarily as a binary classification problem, distinguishing between positive reviews (rating ≥ 4) and lower-rated reviews. For this classification evaluation, we trained a Logistic Regression model using the same text features.

MSE And RMSE allowed us to measure how close each text-based model was to correctly estimating the user's satisfaction based solely on the written review content. AUC was helpful for determining how effectively each text representation captured sentiment and opinion polarity, which are important components of rating prediction.

Textual Feature Evaluation Results

Text Feature	MSE	RMSE	AUC
TF-IDF	1.0356019584749285	1.017645300915269	0.9974759607796853
Bag-of-Words	1.365269002778104	1.1684472614449075	0.9827002205201752
Word2Vec embeddings	1.5019682777765193	1.2255481540015143	0.9316177738597987



Textual Feature Evaluation Results

Based on these results, we adopted TF-IDF as the primary text-representation method in our final hybrid model. TF-IDF performed best because reviews in this dataset are short, highly descriptive, and often rely on domain-specific adjectives and phrases related to fit and comfort. These distinctive terms are weighted strongly by TF-IDF but tend to be overwhelmed in Bag-of-Words and poorly represented in Word2Vec due to the limited context available for training embeddings. The superiority of TF-IDF aligns with the linguistic structure of the dataset and the presence of repeated but meaningful phrases across user reviews. This choice ultimately contributed to the significant improvement observed in the final model.

Temporal Feature Engineering

To incorporate temporal dynamics into the model, we generated several features based on the review date and each item's historical activity. For every review, we calculated the number of days that had passed since it was written, providing a measure of its recency. For each item, we identified the date of its most recent review and derived the number of days since that event, capturing whether the item was actively used or had become outdated. We also computed item popularity based on total review count and extracted the most recent rating each item received to reflect short-term trends in product satisfaction. Finally, the month and year were included to model seasonality, which is especially important in clothing data where demand patterns vary across seasons. Together, these temporal features help the model distinguish between older and more current behavior, detect emerging patterns in item quality, and account for seasonal effects that influence user ratings.

Temporal Feature Engineering

Temporal features improved our model's predictive performance because user behavior and item quality are not static over time. Clothing items often experience shifting popularity, changes in fit perception, and evolving user expectations, all of which are reflected in recent reviews more strongly than older ones. By incorporating features such as days since the review, the time since an item's last recorded review, and the item's most recent rating, the model can identify items that are currently performing well versus those whose quality or fit issues have emerged only recently. Seasonal indicators like month and year also help capture recurring demand cycles related to events, weather, and fashion trends. Without these temporal signals, the model treats older feedback as equally relevant, which can obscure short-term patterns and reduce prediction accuracy.



03

Modeling

Global Mean Baseline



The global mean baseline predicts the same value, 4.55 stars, for every user–item pair, reflecting the strong positivity bias in the RentTheRunway dataset. Its performance (RMSE = 0.716, MAE = 0.588) shows that simply guessing the overall average already produces reasonable accuracy due to the heavily skewed rating distribution. However, because this approach ignores user preferences and item-level differences, it provides only a trivial benchmark. Any useful personalization model must outperform this baseline by leveraging additional user and item information.

User-Item Bias Baseline



The user-item bias model adjusts the global mean by incorporating how much each user typically rates above or below average and how each item is generally perceived. This yields a meaningful improvement over the global mean, achieving an RMSE of 0.787 and an MAE of 0.579. These results show that accounting for user leniency and item popularity captures much of the structure in the rating data. Although simple, this baseline is a strong personalized predictor and provides a solid benchmark that more advanced models must surpass to demonstrate real value.

Machine Learning Problem Formulation

Inputs:

The models utilize a comprehensive set of features to predict star ratings. These can be broadly categorized as: User Metadata, Item Metadata, Temporal Features, Textual Features

Outputs:

The primary output is the stars rating, which is a continuous numerical value ranging from 1 to 5 (derived from the original 2-10 scale).

Optimization:

The task is formulated as a regression problem. The models are optimized to minimize the error between the predicted star ratings and the actual star ratings. Common metrics used for optimization and evaluation are:

- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- Mean Absolute Error (MAE)

Task Type:

This is a supervised regression task, where the goal is to predict a continuous target variable (stars) based on a set of input features.

Model 1: Linear Regression

Linear Regression is a fundamental statistical model that assumes a linear relationship between the input features (X) and the continuous target variable (y).

Appropriateness: Chosen as a strong baseline for structured data. It's appropriate for understanding the direct linear impact of various metadata and temporal features on star ratings.

Advantages:

Simplicity and Interpretability: Easy to understand, implement, and interpret the coefficients, which indicate the magnitude and direction of each feature impact on the rating.

Efficiency: Computationally inexpensive and fast to train, especially with many features.

Good Baseline: Provides a benchmark against which more complex models can be compared.

Disadvantages:

- Assumes Linearity
- Sensitive to Outliers: Outliers can significantly influence the regression line and model coefficients.
- Limited Predictive Power: Often less accurate than more advanced models when dealing with complex datasets with intricate feature interactions.

Complexity/Efficiency/Challenges: Low complexity and high efficiency.

Model 1: Linear Regression

The Linear Regression model, trained with metadata and temporal features, provided the following results:

MSE: 0.5011

RMSE: 0.7079

MAE: 0.5722

These results show a meaningful improvement over the global mean baseline (RMSE of 0.716, MAE of 0.588). The model performs comparably to the user-item bias model (RMSE of 0.787, MAE of 0.579). The RMSE of 0.7079 indicates that, on average, the model's predictions are off by about 0.71 stars, and the MAE of 0.5722 suggests an average absolute prediction error of about 0.57 stars.

This indicates that incorporating user-specific metadata (like height, weight, age, band size, size, and fit code) and temporal features (like days since review, item popularity, item days since last review, item recent rating, month, and year) helps in making more informed rating predictions. Even without including the textual review data, these structured features capture significant patterns and biases in the dataset.



Model 2: Gradient Boosted Regressor



Appropriateness: Selected to capture non-linear relationships and complex interactions within the structured metadata and temporal features. It's well-suited for high-performance prediction tasks.

Advantages:

- High Accuracy: Often achieves state-of-the-art performance on tabular data by effectively handling non-linearities and feature interactions.
- Handles Mixed Data Types
- Feature Importance: Provides a measure of feature importance, indicating which features contribute most to the predictions.

Disadvantages:

- Prone to Overfitting
- Computationally Intensive: Can be slower to train than linear models, especially with large datasets and many trees.

Complexity/Efficiency/Challenges: Moderate complexity, moderate-to-low efficiency. Challenges include hyperparameter tuning to prevent overfitting and managing longer training times.


Model 2: Gradient Boosted Regressor

The Gradient Boosting Regressor, utilizing metadata and temporal features, achieved an **RMSE of 0.6744** and an **MAE of 0.5277**. These metrics signify that the model's predictions are, on average, within approximately 0.67 stars of the actual rating, and the absolute average error is about 0.53 stars.

This performance is notably better than simpler Linear Regression models and baselines, indicating that Gradient Boosting effectively captures non-linear relationships within the structured features, leading to improved star rating predictions.



Model 2: Gradient Boosted Regressor



The graph highlights the most influential features used by the Gradient Boosting Regressor to predict star ratings. The most dominant feature is **fit_code**, which alone accounts for a significant portion of the predictive power (around 46.6% importance). This underscores the critical finding that whether an item fits properly is the primary determinant of a user's rating.

Model 3: Ridge Regression

Ridge Regression is a type of linear regression that includes a regularization term (L2 regularization) to prevent overfitting. When combined with TF-IDF features, it can handle high-dimensional, sparse textual data effectively.

Appropriateness: Chosen to leverage the rich information in the textual reviews (via TF-IDF) in conjunction with structured metadata and temporal features. The regularization is crucial for handling the high dimensionality introduced by TF-IDF, making it suitable for a combined feature set.

Advantages:

- Handles High Dimensionality
- Regularization: Prevents overfitting by shrinking coefficients, leading to better generalization on unseen data.
- Combines Diverse Feature Types: Successfully integrates numerical, categorical, and high-dimensional textual features into a single model.

Disadvantages:

- While robust to high dimensionality, its core is still linear, meaning it might not capture highly complex non-linear interactions as effectively as tree-based methods like Gradient Boosting
- Interpretability of Text Features

Complexity/Efficiency/Challenges: Moderate complexity, high efficiency once features are extracted. Challenges involve managing the feature extraction process for TF-IDF and tuning the regularization parameter (α).

Model 3: Ridge Regression

The Ridge Regression model, which incorporates metadata, temporal features, and TF-IDF for review text, achieved significantly improved performance. It yielded an **MSE of 0.2955**, an **RMSE of 0.5436**, and an **MAE of 0.4124**. These metrics are substantially better than those from the global mean, user-item bias, and text-only (Bag-of-Words, TF-IDF, Word2Vec) models.

This indicates that combining diverse features—user body measurements, item characteristics, temporal dynamics, and detailed textual reviews—is highly effective in predicting star ratings more accurately. The model's predictions are, on average, within half a star of the actual ratings, highlighting the value of a comprehensive feature set for personalized recommendations in this dataset.



04

Evaluation

Evaluation Strategy

The rating prediction task is evaluated using continuous error metrics because the model outputs numerical star ratings rather than discrete categories. The primary metrics used are **Mean Squared Error (MSE)** and **Root Mean Squared Error (RMSE)**, which measure the magnitude of the difference between predicted and true ratings. RMSE is especially appropriate because it penalizes large errors more heavily, which is important in a setting where predicting a rating of five stars instead of one star represents a significantly worse mistake than confusing adjacent values. **Mean Absolute Error (MAE)** is also used to provide a more interpretable measure of average deviation. In addition to these regression metrics, the **text-only models** are evaluated using the **Area Under the ROC Curve (AUC)** by framing the problem as a binary distinction between high-rated and low-rated reviews. This metric assesses how well the text features capture sentiment and is useful for comparing different textual representations. Together, these metrics offer a well-rounded evaluation framework that reflects both predictive accuracy and the ability to capture meaningful patterns in user reviews.

Baselines

Several baselines are included to contextualize the performance of more advanced models. A global mean predictor serves as the simplest baseline and highlights the strong positivity bias present in the dataset. A user-item bias model provides a more informative comparison by accounting for systematic leniency or strictness of individual users and popularity differences across items. Text-only models using Bag-of-Words, TF-IDF, and Word2Vec provide additional baselines focused solely on linguistic information.

Demonstrating improvement over these baselines requires showing that the chosen models consistently produce lower RMSE and MSE, as well as higher AUC for sentiment-based evaluations. The final hybrid Ridge Regression model, which incorporates metadata, temporal features, and TF-IDF text vectors, achieves significantly lower error than any of the baselines. This performance gap demonstrates that combining structured features with text and temporal information yields far more accurate predictions than relying on global averages, user or item biases, or isolated textual representations.

Baseline Results

Baseline	RMSE	MAE	Description
Global Mean Baseline	0.716	0.588	Predicts the mean~4.55 stars always
User Item Bias Baseline	0.787	0.579	Incorporates user and item biases into the mean





Modelling Results

RMSE Improvements (lower is better)

Model	Features Used	RMSE	MAE
Linear Regression	Metadata + temporal	~0.7079	~0.5722
Gradient Boosting	Non-linear structured features	~0.6744	~0.5277
Ridge Regression	Metadata + Temporal + TF-IDF Features	0.5436 ★	~0.4124 ★

Model Summary



The Ridge Hybrid model produces the most accurate predictions. Its error distribution is the narrowest and most centered around zero, indicating fewer and smaller prediction errors compared to the Linear Regression and Gradient Boosting models.

Why our Ridge Regression model is better?

The Ridge Regression model outperforms all other baselines because it is specifically designed to handle the high-dimensional and highly correlated feature space created by combining metadata, temporal signals, and TF-IDF text vectors. TF-IDF alone produces tens of thousands of sparse features, many of which are strongly correlated due to repeated phrases and similar word usage across reviews. Ordinary linear regression cannot reliably estimate coefficients in this setting and becomes unstable, leading to higher variance and larger prediction errors. Ridge Regression addresses this issue through L2 regularization, shrinking coefficients in a way that reduces overfitting while still preserving the full expressive power of the feature set. When temporal features and item-level metadata are added, Ridge is able to integrate these heterogeneous inputs more effectively than models without regularization. Its performance advantage is evident from the reduction in RMSE to 0.5436, significantly lower than the text-only models and non-regularized baselines. These results show that Ridge Regression provides the best balance between flexibility and stability, making it the most reliable model for rating prediction in this dataset.

This can be seen by the reduction of 0.1754 in **RMSE** and 0.1756 in **MAE** compared to the global baseline



05

Discussion of Related Work

Discussion of Related Work



Decomposing fit semantics for product size recommendation in metric spaces
Rishabh Misra, Mengting Wan, Julian McAuley
RecSys, 2018

The RentTheRunway dataset has been used primarily for **fit prediction and size recommendation**, most notably in the RecSys 2018 paper “*Decomposing Fit Semantics for Product Size Recommendation in Metric Spaces*” by Rishabh Mishra, Mengting Wan, and Julian McAuley. The authors introduce the RentTheRunway dataset and show that customers’ self-reported fit labels (“small,” “fit,” “large”) carry rich semantic information about product shape, style, and sizing. They propose latent-factor models and a metric-learning framework to capture these subtle preferences and handle extreme class imbalance inherent in fit feedback .

Other related work on fashion data has focused on **recovering true product sizes**, learning **latent user and item representations**, and modeling **ordinal categories** using logistic or probit regression. Some studies also apply skip-gram embedding models or collaborative filtering to clothing recommendation, but none focus on predicting numeric satisfaction ratings.

Discussion of Related Work

How has prior work approached the same (or similar) tasks?

Prior work treats the RentTheRunway problem as an ordinal classification task, aiming to predict whether an item will fit a user based on their measurements, the item's characteristics, and past transactions. Approaches include:

- Latent-variable models for customer size preferences
- Ordinal regression to preserve the ordering of fit labels
- Metric learning (LMNN) with prototyping to address the uneven distribution of fit labels (see Fig. 1 and Algorithm 1 in RecSys paper)
- Sampling strategies to combat sparsity and imbalance

These methods evaluate performance using AUC, as the task is a three-way classification problem rather than regression.

Discussion of Related Work

How do our results match or differ from what has been reported?

Our study differs fundamentally in its objective: instead of predicting fit, we predict star ratings (1–5), making our task a continuous regression problem. However, we incorporate many of the same fit-related signals emphasized in prior work. In fact, consistent with the findings of Mishra et al., our models show that fit information is extremely predictive.


We use a `fit_code` variable as part of our metadata, and this feature emerges as one of the strongest predictors in our Gradient Boosting Regressor. This aligns with the RecSys 2018 conclusion that customer-reported fit captures important semantic information about how clothing interacts with different body shapes and styles.

While Mishra et al. achieve $AUC \approx 0.72$ for fit classification on RentTheRunway [OBJ], our evaluation uses RMSE and MAE since we predict continuous star ratings. Nonetheless, our results parallel theirs in showing that fit feedback is one of the strongest available predictors, regardless of whether the goal is predicting satisfaction (our work) or predicting size outcomes (their work).

Extension of Prior Research

Our study extends previous research in several ways:

1. **Different Prediction Target:** We shift from predicting fit labels to predicting numerical satisfaction ratings, a task not explored in RecSys 2018.
2. **Hybrid Modeling:** We integrate text (TF-IDF), metadata, and temporal signals—whereas prior work uses only structured transactional attributes.
3. **Temporal Dynamics:** We enhance the feature space using recency, item popularity, and most recent ratings, which were not considered in previous fit-prediction studies.
4. **Demonstrating Fit's Influence on Satisfaction:** Our results empirically show that poor fit strongly predicts lower star ratings, reinforcing the prior paper's claim that fit semantics are central to clothing recommendation.



Thanku you!

